

Solving the Game Content Problem

Pipeline theory and practice

Koray Hagen

Background

Senior Programmer

Tools and Core Technology

SIEA, Santa Monica Studio

Focus:

Animation systems

Content build system

Content compiler tools



Game industry has a big problem

The problem is data

- The time required to develop games is **increasing**

- The time required to develop games is increasing
- The manpower and computational cost of developing games is increasing

- The time required to develop games is increasing
- The manpower and computational cost of developing games is increasing
- The unit price of a game has **remained** relatively stable

Meanwhile ...

- Player expectations for graphical fidelity are **increasing**



- Player expectations for graphical fidelity are increasing
- Player expectations for the scale of the user experience are **increasing**



- Player expectations for graphical fidelity are increasing
- Player expectations for the scale of the user experience are increasing
- Developer expectations for our tools to meet this demand are **increasing**



- This problem has manifested itself in the form of impacted iteration times:
 - For Artists
 - For Designers
 - For Programmers
 - For Production
 - For Operations

- This problem has manifested itself in the form of impacted iteration times:
- Questions:
 - How long until a content author sees the result of changing ____ ?
 - How long until a content author can commit a change to ____ ?

- This problem has manifested itself in the form of impacted iteration times:
- Questions:
- Answers:
 - Desired? Instantly
 - Reality? Possibly minutes to hours

This problem is the **domain** of this lecture

It is also the domain of your game's pipeline

The game content pipeline

- Purpose:
 - Transform unmeaningful data into meaningful data

The game content pipeline

- Purpose:
 - Transform unmeaningful data into meaningful data
 - Often, that means source data into production data

The game content pipeline

- Purpose:
 - Transform unmeaningful data into meaningful data
 - Often, that means source data into production data
 - Production data is content that ends up in a player's hands

The game content pipeline

- Purpose:
- Game content:
 - Meshes
 - Textures
 - Animations
 - Shaders
 - Materials
 - Physics
 - Scripts
 - Narrative
 - Audio

The game content pipeline

- Purpose:
- Game content:
 - Heterogeneous:
 - Has different intrinsic characteristics (formats, disk-space impact)
 - Requires different applied transformations

The game content pipeline

- Purpose:
- Game content:
 - Heterogeneous:
 - Has different intrinsic characteristics (formats, disk-space impact)
 - Requires different applied transformations
 - Hierarchical relationships
 - Dependencies
 - Content identity

The game content pipeline

- Philosophy:
 - Create a **unified** content pipeline system and architecture at our studio

The game content pipeline

- Philosophy:
 - Create a unified content pipeline system and architecture at our studio
 - Properties:
 - Correct
 - Expressive
 - Explicit
 - Homogenous
 - Scalable
 - Debuggable

This is a lecture on how we **think** and **reason** about achieving these desired properties

The agenda


- Foundations
 - Creating the building blocks of a content pipeline

The agenda

- Foundations
- Theory and architecture
 - Considerations for building a pipeline
 - High level view of design

The agenda

- Foundations
- Theory and architecture
- Lessons and current research
 - Concrete technology examples
 - Questions for the future



Foundations

Each piece represents a **conceptual** building block

Data transformation functions

- Purpose:
 - Core operator of a game's pipeline
 - Transforms a **unit** of data from one state to another state

Data transformation functions

- Purpose:
- Example:

```
1 struct mesh { /* members */ };
2 struct mesh_opt { /* members */ };
3
4 bool mesh_to_optimized_mesh(const mesh* src, mesh_opt* dst);
5 bool merge_optimized_mesh(const mesh_opt** src, mesh_opt* dst);
```

Data transformation functions

- Purpose:
- Example:
- Invariants:
 - Serializable
 - Transferrable
 - Side-effect free
 - Discrete
 - Atomic

Data transformation functions

- Purpose:
- Example:
- Invariants:
 - Serializable
 - Both inputs and outputs are representable in different contexts
 - Memory mapped
 - File system
 - Networked

Data transformation functions

- Purpose:
- Example:
- Invariants:
 - Serializable
 - Transferrable
 - Usable within a game's pipeline or run-time

Data transformation functions

- Purpose:
- Example:
- Invariants:
 - Serializable
 - Transferrable
 - Side-effect free
 - Safely usable in a concurrent system
 - Knowable state and always **debuggable**

Data transformation functions

- Purpose:
- Example:
- Invariants:
 - Serializable
 - Transferrable
 - Side-effect free
 - Discrete
 - Conceptually simple to reason about
 - Represents a single **unit** of transformation

Data transformation functions

- Purpose:
- Example:
- Invariants:
 - Serializable
 - Transferrable
 - Side-effect free
 - Discrete
 - Atomic
 - Transformation either **did** or **did not** occur, there is no half-way

Data transformation functions

- Goals:
 - Develop a library of transformations
 - Define an appropriate granularity
 - Feverously respect invariants

Data transformation functions

- Goals:
 - Develop a library of transformations
 - Needed to represent the full scope of game data
 - Define an appropriate granularity
 - Feverously respect invariants

Data transformation functions


- Goals:
 - Develop a library of transformations
 - Define an appropriate granularity
 - Base judgment on time and resources required for the transform
 - Feverously respect invariants

Data transformation functions

- Goals:
 - Develop a library of transformations
 - Define an appropriate granularity
 - Feverously respect invariants
 - This will allow powerful architectural opportunities later

Dependency graph

- Purpose:
 - Track the global state of data and **relationships**
 - Serves as an interface to both data and transformations



Dependency graph

- Purpose:
- Data dependency graph
- Transformation dependency graph

Dependency graph

- Purpose:
- Data dependency graph
 - The **relationship** network of data files
 - Example:
 - Skybox
 - Cubemap
 - 1D Texture (six)
 - Cubemap Texture
- Transformation dependency graph

Dependency graph

- Purpose:
- Data dependency graph
 - The relationship network of data files
 - Example:
 - Properties:
 - File identity (as a path or GUID)
 - Connectivity to and from other files
 - Reachability to and from other files
- Transformation dependency graph

Dependency graph

- Purpose:
- Data dependency graph
- Transformation dependency graph
 - The relationship network of data transformation functions
 - Example:
 - `bool cubemap_to_skybox(const cubemap* src, skybox* dst);`
 - `bool textures_to_cubemap(const texture** src, cubemap* dst);`
 - `bool tga_to_texture(const tga* src, texture* dst);`

Dependency graph

- Purpose:
- Data dependency graph
- Transformation dependency graph
 - The relationship network of data transformation functions
 - Example:
 - Properties and differences from data dependency graph:
 - Requires **critical path** evaluation for a desired result
 - Requires **deciding** if and when a transformation occurs
 - Performance critical

Storage

- Purpose:
 - Provide permanent or transient contexts for information and data



Storage

- Purpose:
- Considerations:
 - Often a tradeoff between **latency**, **capacity**, and **longevity**
 - Example: low latency, in-memory key-value database
 - Example: high latency, networked file system

Storage

- Purpose:
- Considerations:
 - Often a tradeoff between latency, capacity, and longevity
 - Needed for many areas in a pipeline:
 - Source data
 - Intermediate data
 - Inter/Intraprocess communication
 - Production data

Storage

- Purpose:
- Considerations:
 - Often a tradeoff between latency, capacity, and longevity
 - Needed for many areas in a pipeline:
 - Should never be directly interacted with by **transformation functions**:
 - Fundamentally breaks invariants
 - Should be orchestrated at a higher level
 - During evaluation of the **transformation dependency graph**



Theory and architecture

Computational work

- Iteration times:
 - For content authors, improving iteration is purely a metric of time
 - The amount of computation required is not relevant

Computational work

- Iteration times:
 - For content authors, improving iteration is purely a metric of time
 - The amount of computation required is not relevant
 - Often, more computation results in more time required

Computational work

- Iteration times:
- Mitigation strategies:
 - Data model
 - Concurrency
 - Avoidance

Computational work

- Iteration times:
- Mitigation strategies:
 - Data model
 - Reduction of transformation steps required to see result
 - Transformations are unnecessary or optional
 - Direct optimization or removal of transformations
 - Concurrency
 - Avoidance

Computational work

- Iteration times:
- Mitigation strategies:
 - Data model
 - Reduction of transformation steps required to see result
 - Ideal to address, the best general optimization is to do **less work**
 - Concurrency
 - Avoidance

Computational work

- Iteration times:
- Mitigation strategies:
 - Data model
 - Concurrency
 - Leveraging the **invariants** of our data transformation functions
 - Available contexts:
 - Threads
 - Processes
 - Network
 - Avoidance

Computational work

- Iteration times:
- Mitigation strategies:
 - Data model
 - Concurrency
 - Avoidance
 - Hierarchical caching contexts (more detail later):
 - In-memory
 - File system
 - Network

Computational work

- Iteration times:
- Mitigation strategies:
 - Data model
 - Concurrency
 - Avoidance
 - Hierarchical caching contexts:
 - Pruning action based on criteria:
 - Unneeded for desired result
 - Redundant

Conceptual architectures

- Live iteration
- Asset baking

Conceptual architectures

- Live iteration
 - Working directly in the game to see changes
 - Streaming data to the game through an external tool
- Asset baking

Conceptual architectures

- Live iteration
 - Working directly in the game to see changes
 - Streaming data to the game through an external tool
 - Properties:
 - Has the benefit of user seeing **instant** results
 - Has the problem of data **transience**
 - Edit state could become corrupt
 - Edit state could encounter an irrecoverable error (crashes)
 - Has the problem of needing to reflect changes into source data
 - Other assets may have **dependence** on changes
- Asset baking

Conceptual architectures

- Live iteration
- Asset baking
 - Using an offline toolset to “cook” or “bake” assets for the game

Conceptual architectures

- Live iteration
- Asset baking
 - Using an offline toolset to “cook” or “bake” assets for the game
 - Properties:
 - Has the benefit of **permanence**
 - Has the benefit of simple directional flow of changes
 - Source → intermediate → production
 - Has the problem, often, of **increased computational cost**
 - Has the problem, often, of **non-uniformity**
 - Example: Different teams having different file formats
 - Resulting in a more **complex** unified pipeline

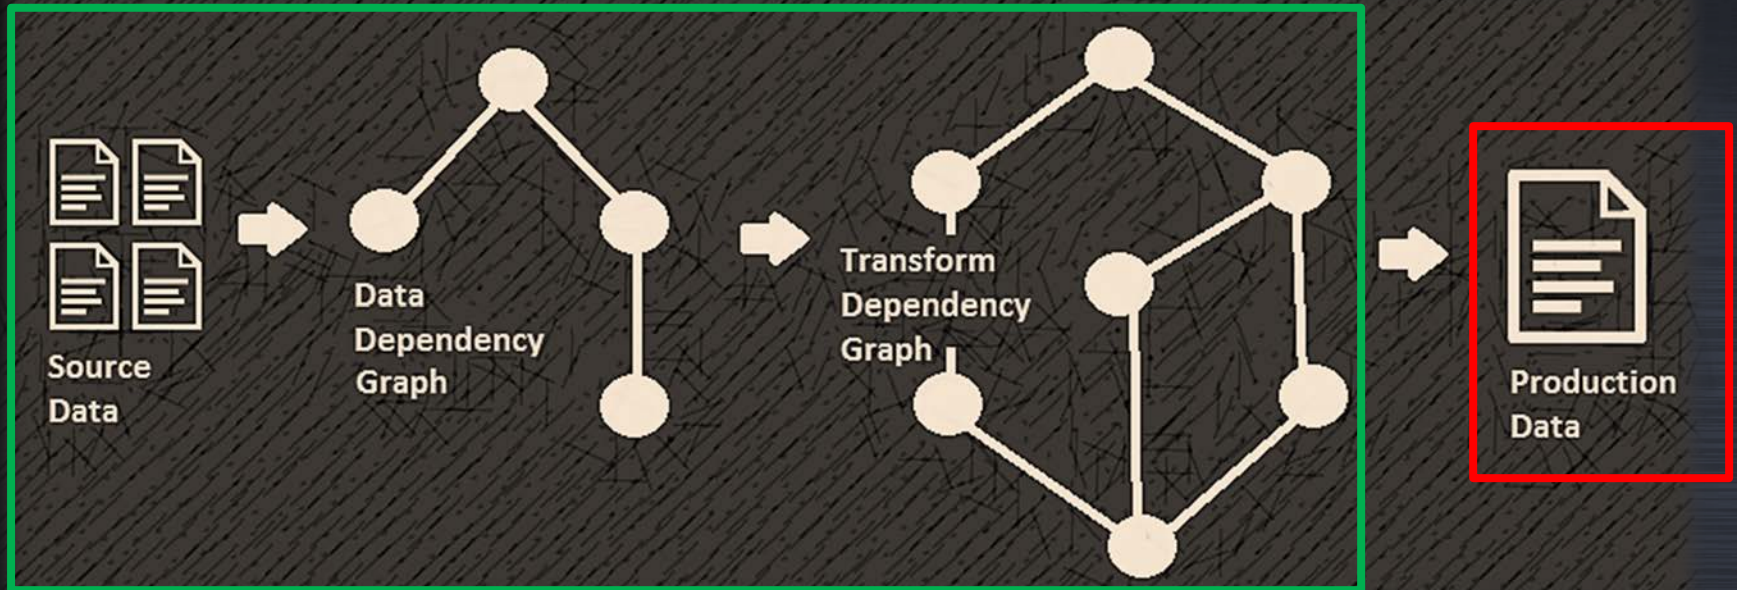
Conceptual architectures

- Live iteration
- Asset baking
- Philosophy:
 - Live iteration workflows and asset baking can be **reconciled**
 - Both concepts are **shades** of the same architectural principles

Conceptual architectures

- Live iteration
- Asset baking
- Philosophy:
- Desired pipeline architecture:
 - **Transferrable** library of available data transformations
 - Full data dependency graph of all game content
 - **Computable** transformation dependency graphs
 - **Generalized** hierarchical caching model

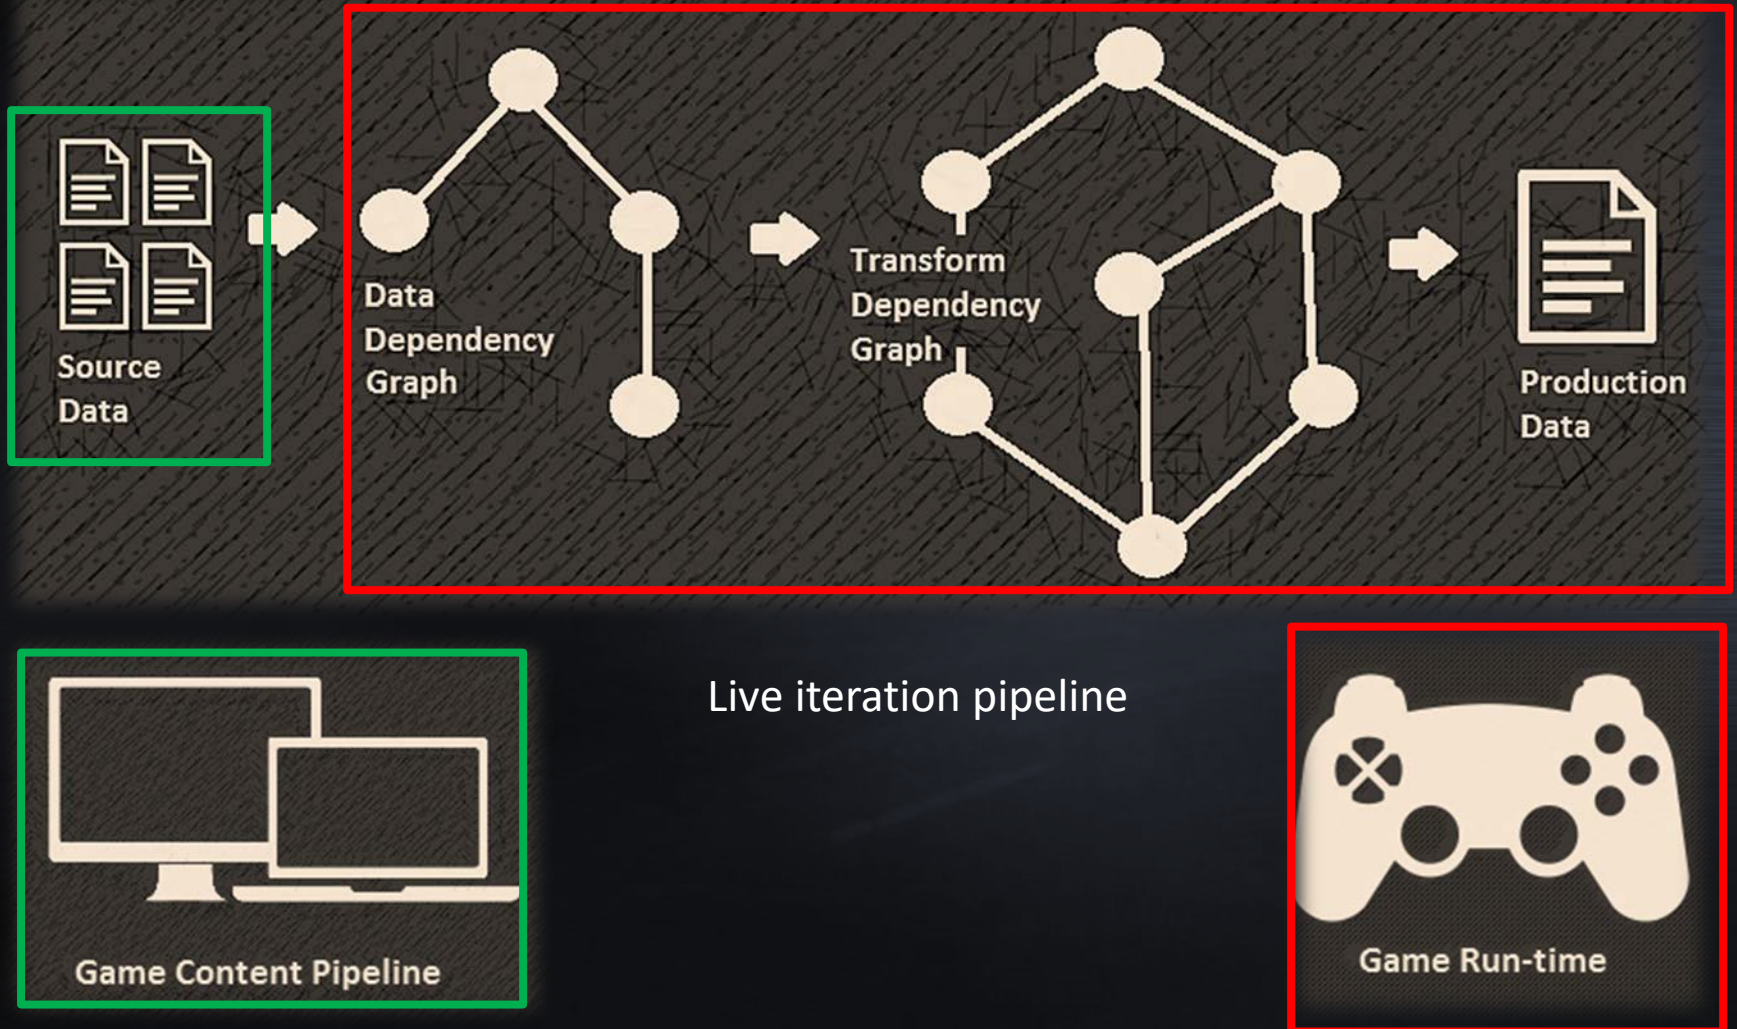
Visualization of a content build




Asset baking pipeline



Visualization of a content build





Lessons and current research

Data transformation functions

- Lesson: Define the data formats
- Lesson: Discipline in respecting invariants
- Lesson: Transferability is crucial

Data transformation functions

- Lesson: Define the data formats
 - Developed an in-house object serialization library
 - Similar to [Google Flat Buffers](#) and [Protocol Buffers](#)
- Lesson: Discipline in respecting invariants
- Lesson: Transferability is crucial

Data transformation functions

- Lesson: Define the data formats
 - Developed an in-house object serialization library
 - Similar to Google Flat Buffers and Protocol Buffers
 - Designed for:
 - Memory-mappability
 - Binary format for **speed**, Json format for **debugging**
 - Data layout optimization
 - Versioning and Identification
 - Automatic API generation
 - Custom allocators
- Lesson: Discipline in respecting invariants
- Lesson: Transferability is crucial

Data transformation functions

- Lesson: Define the data formats
- Lesson: Discipline in respecting invariants
 - Allows for **concurrency** guarantees
 - Allows for data, build, and game **state** guarantees
- Lesson: Transferability is crucial

Data transformation functions

- Lesson: Define the data formats
- Lesson: Discipline in respecting invariants
- Lesson: Transferability is crucial
 - Same library of functions can be used at **bake time** and **run time**
 - Opportunity for the transformation dependency graph to **computationally decide** where to execute transform:
 - A thread
 - A local process
 - A networked process

Dependency Graphs

- Lesson: Most research required here
- Lesson: Describing the flow of a network
- Lesson: Graph operators

Dependency Graphs

- Lesson: Most research required here
 - *Open problem*: Scaling computation to **millions of nodes**
 - *Open problem*: Cache coherent spatial data structures and processing
 - *Open Problem*: Extremely performant transformation dependency graph evaluation and **iterative construction**
- Lesson: Describing the flow of a network
- Lesson: Graph operators

Dependency Graphs

- Lesson: Most research required here
- Lesson: Describing the flow of a network
 - Transformation dependency graph
 - Required a language (**DSL**) to describe the transformation network
 - Transformations known as **commands**:
 - Examples of granularity:
 - » **compress_texture**
 - » **create_animation_clip**
 - DSL results in **vast** amount of automatic code generation
- Lesson: Graph operators

Dependency Graphs

- Lesson: Most research required here
- Lesson: Describing the flow of a network
- Lesson: Graph operators
 - Transformation dependency graph traversal:
 - Two **operators** could encapsulate our most **complex** operations
 - Reductions:
 - Compute a **subgraph** under a node as input to a command
 - Maps:
 - Direct input node to a command

Storage


- Lesson: File system by itself is very limiting
- Lesson: Creating a global asset library

Storage

- Lesson: File system by itself is very limiting
 - Butting heads with **Windows** file system rules is painful:
 - Path lengths (256 characters)
 - Path and file name **uniqueness**
 - Often results in **insane** naming conventions for files
- Lesson: Creating a global asset library

Storage

- Lesson: File system by itself is very limiting
- Lesson: Creating a global asset library
 - Many reasons for desirability:
 - Freedom for custom **identification** and file **versioning**
 - Freedom from file system path and naming **conventions**
 - Filenames on physical disk can be **mangled**, while the asset library can ensure **human readability**



Caching

- Lesson: Hierarchical caching is a powerful model
 - Goal: **Never re-compute** a function with the same parameters.
 - Achieved by decorating functions with **appropriate** caching strategies:
 - Match a function with the most appropriate caching semantic
 - Driven by the nature of input and output data
 - Example:
 - File-based parameters cached to the file system or network
 - Normal function parameters cached in memory

Caching

- Lesson: Hierarchical caching is a **powerful** model
 - Any piece of computation can be **decorated** with caching
 - Found it to be most beneficial with:
 - Path and string manipulation
 - Data transformation functions

Caching

- Lesson: Hierarchical caching is a powerful model
 - Any piece of computation can be decorated with caching
 - Hiding caching structure behind generalized **interfaces** and **decorators** allowed for experimentation.
 - Proved to be invaluable when we switched our file-system based key-value store from **SQLite** to **LMDB**

Caching

- Lesson: Hierarchical caching is a powerful model
 - Any piece of computation can be decorated with caching
 - Hiding caching structure behind generalized interfaces and decorators allowed for experimentation.
 - Our current implementations:
 - In-memory: achieved with function level **memoization**
 - File-system: achieved with **LMDB**
 - Networked: achieved with an in-house technology backed by **CEPH** (a networked file system object store)

Caching

- Lesson: Hierarchical caching is a **powerful** model

- Example:

```
1 [global_cached]
2 void execute_command(command* func, const file_node* src, file_node* dst) {
3     // func merge_optimized_mesh
4     func(src.get(), dst.get())
5 }
6
7 [memoized]
8 bool expensive_operation(mesh* src) {
9     // compute something expensive
10    return true;
11 }
12
13 bool merge_optimized_mesh(const mesh_opt** src, mesh_opt* dst) {
14    return expensive_operation(dst)
15 }
```



Takeaways

Hard Problems

- Building the foundation
- Dependency graph scalability
- Dependency graph decision making

Hard Problems

- Building the foundation
 - Most studios, including ours have a **long history** of tool's development
 - Difficult to modernize and **justify** changing monolithic, **stable** tools
 - Large upfront cost for changing data formats
- Dependency graph scalability
- Dependency graph decision making

Hard Problems

- Building the foundation
- Dependency graph scalability
- Dependency graph decision making
 - Performant and meaningful computation against large graphs is **difficult** but not **insurmountable**.
 - Pushing automated decision making to cover:
 - Caching strategies
 - Appropriate contexts for data transformations

The problem is data

The solution is automation

Thank you
Questions?